

Deep Learning for Computer Vision

Ahmed Hosny Abdel-Gawad

Senior AI/CV Engineer



Table Content

Part One

From Traditional Computer Vision to Deep Learning

Part Two

Practical Convolution Neural Nets

Part Three

Computer Vision Applications

1 From Traditional CV to DL

1.1

From Traditional Convolution Neural Networks to Deep Learning

Soft introduction about OpenCV and what exactly is traditional computer vision.

1.2

Convolution Neural Networks (CNNs)

Motivations behind CNNs, Image Classifications, calculation of sizes of filters, input and output layers.

1.3

Convolution Neural Network Meta Architectures

How to Design CNNs and the well-known architectures.

1.1 From Traditional ConvNets to DL

Semantic gap and CV tasks

The gap between what a computer sees and what we want it to see.

Traditional CV Pipeline

Traditional images processing techniques and its limitations.


From traditional to learnable convolution filters

Introducing the learning processes and convolution neural networks.

ConvNets and GPU

Why GPUs works perfect with convolution neural networks.

Semantic Gap



What Does A Computer See?

Using RGB model to help us understand how a computer looks at an image.

02

What We Want the Computer to see?


The core computer vision task: Image classification.

03

OpenCV Primer

Introduction to OpenCV.

What Does A Computer See?



This image by Nikita is
licensed under [CC-BY 2.0](#)

```
[[105 112 108 111 104 99 106 99 96 103 112 119 104 97 93 87]
 [ 91 98 102 106 104 70 98 103 99 105 123 136 110 105 94 85]
 [ 76 85 90 105 128 105 87 96 95 99 115 112 106 103 99 85]
 [ 99 81 81 93 120 131 127 100 95 98 102 99 96 93 101 94]
 [106 91 61 64 69 91 88 85 101 107 109 98 75 84 96 95]
 [114 108 85 55 55 69 64 54 64 87 112 129 98 74 84 91]
 [133 137 147 103 65 81 80 65 52 54 74 84 102 93 85 82]
 [128 137 144 140 109 95 86 70 62 65 63 63 60 73 86 101]
 [125 133 148 137 119 121 117 94 65 79 80 65 54 64 72 98]
 [127 125 131 147 133 127 126 131 111 96 89 75 61 64 72 84]
 [115 114 109 123 150 148 131 118 113 109 100 92 74 65 72 78]
 [ 89 93 90 97 108 147 131 118 113 114 113 109 106 95 77 80]
 [ 63 77 86 81 77 79 102 123 117 115 117 125 125 130 115 87]
 [ 62 65 82 89 78 71 80 101 124 126 119 101 107 114 131 119]
 [ 63 65 75 88 89 71 62 81 120 138 135 105 81 98 110 118]
 [ 87 65 71 87 106 95 69 45 76 130 126 107 92 94 105 112]
 [118 97 82 86 117 123 116 66 41 51 05 03 80 05 102 107]
 [164 146 112 80 82 120 124 104 76 48 45 66 88 101 102 109]
 [157 170 157 120 93 86 114 132 112 97 60 55 70 82 99 94]
 [130 128 134 161 139 100 109 118 121 134 114 87 65 53 69 86]
 [120 112 96 117 150 144 120 115 104 107 102 93 87 81 72 79]
 [123 107 96 86 83 112 153 149 122 109 104 75 80 107 112 99]
 [122 121 102 80 82 86 94 117 145 148 153 102 58 78 92 107]
 [122 164 148 103 71 56 78 83 93 103 119 139 102 61 69 84]]
```

What the computer sees



An image is just a big grid of numbers between [0, 255]:

e.g. 800 x 600 x 3
(3 channels RGB)

What Does A Computer See?

Every major color has a range from 0 to 255, we can infer that higher the value, the brighter is the color.


When we combine two colors, say red and green, the resulting color is Yellow. It is represented in the three dimensional space as 255,255,0 (R,G,B).



What Does A Computer See?

An image is made up of pixels placed adjacent to one another. These colored pixels are made up of three channels which are placed one behind the other.

The below pixelated image is a combination of the three channels which are placed one behind the other.



Basics of image processing with **filtering**

When it comes to detecting edges and contours, noise gives a great impact on the accuracy of detection. Therefore removing noises and controlling the intensity of the pixel values can help the model to focus on the general details and get higher accuracy.

=

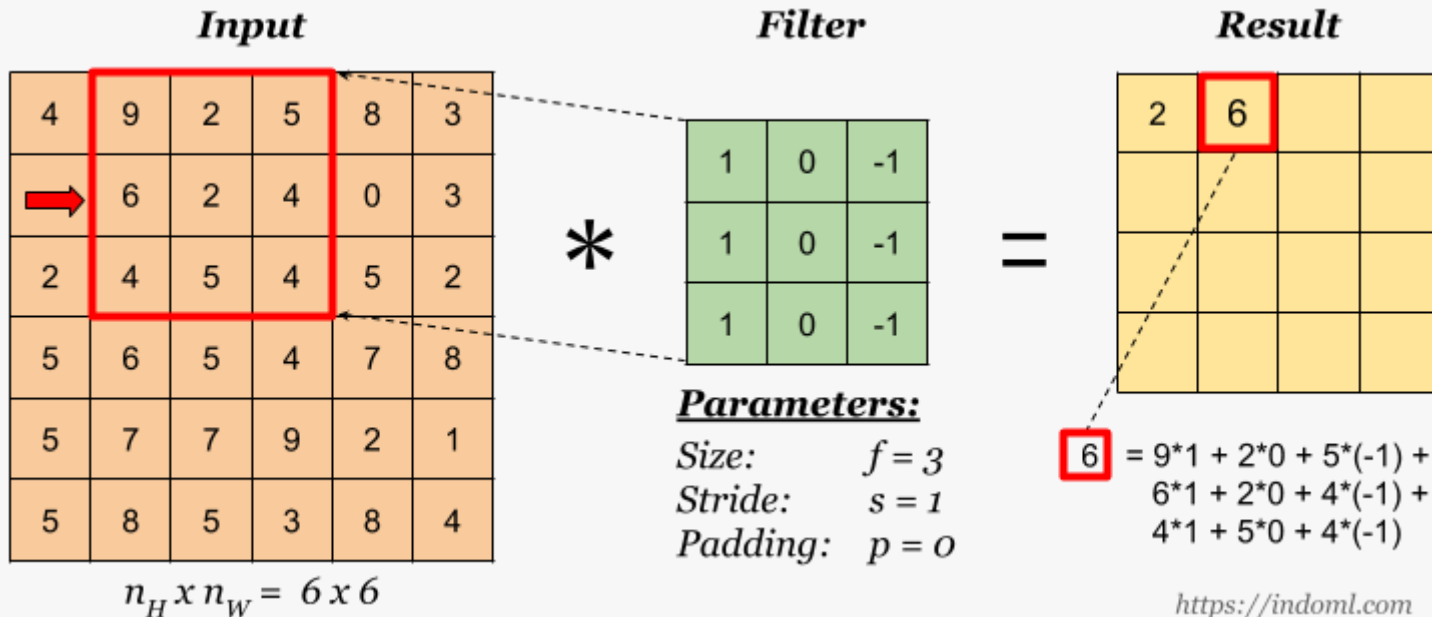
Blurring, thresholding, and morphological transformation are the techniques we use for this purpose.

Blurring

The main goal of blurring is to perform noise reduction.

Four major blurring techniques used in OpenCV: **Averaging blurring**, **Gaussian blurring**, **median blurring** and bilateral filtering.

All four techniques have a common basic principle, which is applying **convolutional operations** to the image with a filter (**kernel**). The values of the applying filters are different between the four blurring methods.



Average Blurring

Average blurring is taking the average of all the pixel values under the given kernel area and replace the value at the center. **Then what will it be like if we increase the size of the kernel?**

$$\frac{1}{25} \begin{bmatrix} 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \end{bmatrix}$$




156	124	112	150	130	142
128	180	130	144	145	140
111	150	110	200	142	128
101	130	170	250	133	120
121	150	200	170	119	150
130	150	150	135	200	150

Gaussian/Bilateral **Blurring**

Gaussian blurring is nothing but using the kernel whose values have a Gaussian distribution. The values are generated by a Gaussian function so it requires a sigma value for its parameter.

Bilateral
noises but
uses Gaussian
Therefore




olving
s. It
count.

The shape of a Gaussian filter (on the left) and a Bilateral filter (on the right)

Thresholding

Thresholding transforms images into binary images.


We need to set the threshold value and max values and then we convert the pixel values accordingly. There are five different types of thresholding: Binary, the inverse of Binary, Threshold to zero, the inverse of Threshold to Zero, and Threshold truncation. [link](#)



Adaptive Thresholding

What if we have a picture with various amount of lighting in different areas?


In this case, applying one value to the whole image would be a bad choice. A better approach would be using different thresholds for each part of the image. **Adaptive thresholding**, which serves this issue. By calculating the threshold within the neighborhood area of the image, we can achieve a better result from images with varying illumination.



Gradient

In mathematics, the gradient geometrically represents the slope of the graph of a function with multi-variables.

The image gradient represents directional changes in the intensity or color mode and we can use this concept for locating edges.



Morphological transformations

It's also possible to manipulate the figures of images by filtering, which is called as **morphological transformation**.

Erosion is the technique for shrinking figures and it's usually processed in a grayscale. The shape of filters can be a **rectangle**, an **ellipse**, and a **cross** shape. By applying a filter we remove any 0 values under the given area.

0	0	0	0	0	0	0
0	0	1	0	0	0	0
0	0	1	1	1	0	0
0	1	1	1	1	0	0
0	1	1	1	1	0	0
0	1	1	1	1	0	0
0	0	0	0	0	0	0

Erosion

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	0	1	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

Morphological transformations

Erosion is the technique for shrinking figures and it's usually processed in a grayscale. The shape of filters can be a **rectangle**, an **ellipse**, and a **cross** shape. By applying a filter we remove any 0 values under the given area.



Original Image



Basic kernel



Ellipse kernel



Cross kernel

Morphological transformations

It's also possible to manipulate the figures of images by filtering, which is called as **morphological transformation**.

Dilation is the opposite to erosion. It is making objects expand and the operation will be also opposite to that of erosion.


0	0	0	0	0	0	0
0	0	\	0	0	0	0
0	0	\	\	\	0	0
0	\	\	\	\	0	0
0	\	\	\	\	0	0
0	\	\	\	\	0	0
0	0	0	0	0	0	0

Dilation

0	0	0	0	0	0	0
0	\	\	\	\	\	0
\	\	\	\	\	\	0
\	\	\	\	\	\	0
\	\	\	\	\	\	0
\	\	\	\	\	\	0
\	\	\	\	\	\	0


Morphological transformations

Dilation is the opposite to erosion. It is making objects expand and the operation will be also opposite to that of erosion.



Morphological transformations

Opening and **closing** operation is the mixed version of erosion and dilation. Opening performs erosion first and then dilation is performed on the result from the erosion while closing performs dilation first and the erosion.



Morphological transformations

Top hat filter is the subtracted area from opening to the original image while **black hat filter** is that from closing.



Original Image



Opening



Closing



Gradient



Tophat



Blackhot

A core task in CV: Image Classification

Image Classification




This image by Nikita is
licensed under [CC-BY 2.0](#)

(assume given set of discrete labels)
{dog, cat, truck, plane, ...}

→ cat

The Problem: Semantic Gap



This image by Nikita is
licensed under [CC-BY 2.0](#)


```
[[105 112 108 111 104 99 106 99 96 103 112 119 104 97 93 87]
 [ 91 98 102 106 104 70 98 103 99 105 123 136 110 105 94 86]
 [ 76 85 90 105 128 105 87 96 95 99 115 112 106 103 99 85]
 [ 99 81 81 93 120 131 127 100 95 98 102 99 96 93 101 94]
 [106 91 61 64 69 91 88 85 101 107 109 98 75 84 96 95]
 [114 108 85 55 55 69 64 54 64 87 112 129 98 74 84 91]
 [133 137 147 103 65 81 80 65 52 54 74 84 102 93 85 82]
 [128 137 144 140 109 95 86 70 62 65 63 63 60 73 86 101]
 [125 133 148 137 119 121 117 94 65 79 80 65 54 64 72 98]
 [127 125 131 147 133 127 126 131 111 96 89 75 61 64 72 84]
 [115 114 109 123 150 148 131 118 113 109 100 92 74 65 72 78]
 [ 89 93 90 97 108 147 131 118 113 114 113 109 106 95 77 80]
 [ 63 77 86 81 77 79 102 123 117 115 117 125 125 130 115 87]
 [ 62 65 82 89 78 71 80 101 124 126 119 101 107 114 131 119]
 [ 63 65 75 88 89 71 62 81 120 138 135 105 81 98 110 118]
 [ 87 65 71 87 106 95 69 45 76 130 126 107 92 94 105 112]
 [118 97 82 86 117 123 116 66 41 51 05 03 80 05 102 107]
 [164 146 112 80 82 120 124 104 76 48 45 66 88 101 102 109]
 [157 170 157 120 93 86 114 132 112 97 60 55 70 82 99 94]
 [130 128 134 161 139 100 109 118 121 134 114 87 65 53 69 86]
 [120 112 96 117 150 144 120 115 104 107 102 93 87 81 72 79]
 [123 107 96 86 83 112 153 149 122 109 104 75 80 107 112 99]
 [122 121 102 80 82 86 94 117 145 148 153 102 58 78 92 107]
 [122 164 148 103 71 56 78 83 93 103 119 139 102 61 69 84]]
```

What the computer sees

An image is just a big grid of numbers between [0, 255]:

e.g. 800 x 600 x 3
(3 channels RGB)

Challenges: Viewpoint variation






All pixels change when
the camera moves!

Challenges: Illumination



We need the computer to
classify all of them as cats!

Challenges: Deformation



Again, we need the computer to classify all of them as cats!

Challenges: Occlusion



Challenges: Background Clutter



Challenges: Intraclass variation



An Image Classifier




```
def classify_image(image):  
    # Some magic?  
    return class_label
```

Unlike e.g. sorting a list of numbers,

No obvious way to hard-code the algorithm for recognizing a cat, or other classes.

Attempts: Traditional CV



Layers of features


Pixels



Gray Pixels




Threshold Pixels







Edges/Dilation



Contours



Traditional LPD



Thank You!